

Monte Carlo Integration

Michael Rose

November 28, 2018

3.1 | Introduction

Two major classes of numerical problems in statistical inference are *optimization* and *integration*. These are in contrast to the analytical methods of computing estimators like maximum likelihood, bayes, method of moments, and others.

A general solution to the numerical solutions needed for analytically intractable integrals is to use simulations, of either the true or some substitute distributions, to calculate the quantities of interest.

Note that the possibility of producing an almost infinite number of random variables distributed according to a given distribution gives us access to the use of frequentist and asymptotic results much more easily than in the usual inferential settings, where the sample size is most often fixed.

Before this, it is worth noting the alternative to monte carlo integration: numerical methods like simpsons and the trapezoidal rule. R also provides functions for unidimensional integration, `area()` which cannot handle infinite bounds, and `integrate()` which can, but is very fragile.

Example 3.1

As a test, lets compare the use of `integrate()` on the integral


$$\int_0^{\infty} x^{\lambda-1} \exp(-x) dx$$

with the computation of $\Gamma(\lambda)$ via the `gamma()` function.

```
# create function to get integrate() values
int_gam <- function(lambda){
  integrate(function(x){
    x^(lambda - 1) * exp(-x)
  }, 0, Inf)$val
}

# make data frames
log_int_gam <- seq(0.01, 10, length.out = 100) %>% sapply(int_gam) %>% log() %>% tibble("output" = .)
log_gam_fun <- seq(0.01, 10, length.out = 100) %>% sapply(lgamma) %>% tibble("output" = .)

# plot
ggplot() + geom_point(aes(x = log_int_gam$output, y = log_gam_fun$output)) + coord_fixed() +
  xlab("log(integrate(f))") + ylab(expression(log(Gamma(lambda))))
```



The figure above shows that there is not a discrepancy, even for small values of lambda.

A main difficulty with numerical integration methods like `integrate()` is that they often fail to spot the region of importance for the function to be integrated. In contrast, simulation methods naturally tend to target this region by exploiting the information provided by the probability density associated with the integrals.

Example 3.2

Consider a sample of ten Cauchy RVs x_i , $1 \leq i \leq 10$ with a location parameter $\theta = 350$. Then the pseudo-marginal distribution under a flat prior is

$$m(x) = \int_{-\infty}^{\infty} \prod_{i=1}^{10} \frac{1}{\pi} \frac{1}{1+(x_i-\theta)^2} d\theta.$$

However, `integrate()` gives the wrong value and fails to signal the difficulty since the error evaluation is absurdly small:

```
# define 10 cauchys
ten_cauchy <- rcauchy(10) + 350

# take the product of ten cauchys
likelihood <- function(theta){
  u <- dcauchy(ten_cauchy[1] - theta)
  for (i in 2:10){
    u <- u * dcauchy(ten_cauchy[i] - theta)
  }
  return(u)
}
```

```

}

# integrate product over +- inf
integrate(likelihood, -Inf, Inf)

## 4.805305e-44 with absolute error < 9.5e-44

Furthermore, the result is not comparable to area:

# define 10 cauchys
ten_cauchy <- rcauchy(10)

# create integrate() function
cauc_int <- function(a){
  integrate(likelihood, -a, a)$val
}

# create area() function
cauc_area <- function(a){
  MASS::area(likelihood, -a, a)
}

# create vector of thetas
thetas <- seq(1, 10^3, length.out = 10^4)

# make dataframes
cauc_integrate <- thetas %>% sapply(cauc_int) %>% log() %>% tibble("output" = .)
cauc_area <- thetas %>% sapply(cauc_area) %>% log() %>% tibble("output" = .)

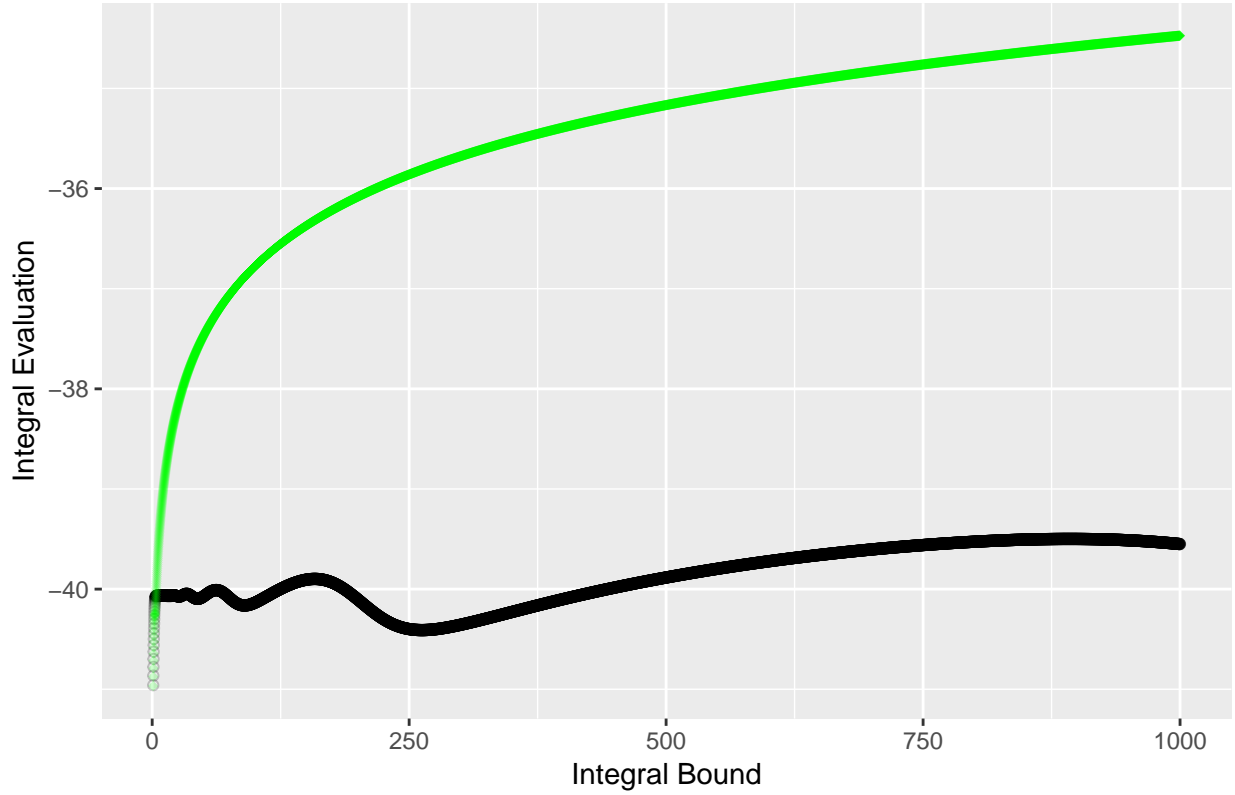
# find dataframe range
bounds <- cbind(cauc_area, cauc_integrate) %>% range()
bounds

## [1] -40.96781 -34.47310

ggplot() +
  geom_point(aes(x = thetas, y = cauc_integrate$output), shape = 1, alpha = 0.2) +
  geom_point(aes(x = thetas, y = cauc_area$output), shape = 18, color = 'green', alpha = 0.2) +
  ylim(bounds) + xlab("Integral Bound") + ylab("Integral Evaluation") +
  ggtitle("Integrate in Black, Area in Green")

```

Integrate in Black, Area in Green



We can see that using `area()` in this case produces a more reliable evaluation, since it flattens out as a increases.

3.2 | Classical Monte Carlo Integration

Suppose we wish to evaluate the integral

$$\mathbb{E}_f[h(X)] = \int_{\chi} h(x)f(x)dx$$

where χ denotes the set where the random variable X takes its values, which is usually equal to the support of the density f .

The principle of the Monte Carlo method is to generate a sample (X_1, \dots, X_n) from the density f and propose as an approximation the empirical average

$$\bar{h}_n = \frac{1}{n} \sum_{j=1}^n h(x_j)$$

computed by `mean(h(x))` in R, since \bar{h}_n converges almost surely (i.e. for almost every generated sequence) to $\mathbb{E}_f[h(X)]$ by the Strong Law of Large Numbers.

When $h^2(X)$ has a finite expectation under f , the speed of convergence of \bar{h}_n can be assessed since the convergence takes place at a rate of $O(\sqrt{n})$ and the asymptotic variance of the approximation is

$$\text{var}(\bar{h}_n) = \frac{1}{n} \int_{\chi} (h(x) - \mathbb{E}_f[h(X)])^2 f(x)dx$$

which can also be estimated from the sample (X_1, \dots, X_n) through

$$v_n = \frac{1}{n^2} \sum_{j=1}^n [h(x_j) - \bar{h}_n]^2$$

More specifically, due to the **Central Limit Theorem**, for large n ,

$$\frac{\bar{h}_n - \mathbb{E}_f[h(X)]}{\sqrt{v_n}} \sim \mathcal{N}(0, 1)$$

which leads to the construction of a convergence test and confidence bounds on the approximation of $\mathbb{E}_f[h(X)]$.

Example 3.3


Given $h(x) = [\cos(50x) + \sin(20x)]^2$, consider evaluating its integral over $[0, 1]$. It can be seen as a uniform expectation, so we can generate U_1, U_2, \dots, U_n iid $\mathcal{U}(0, 1)$ random variables and approximate $\int h(x)dx$ with $\sum \frac{h(U_i)}{n}$.

```
# define function
sin_cos_fn <- function(x){
  return((cos(50 * x) + sin(20 * x))^2)
}

# get integrate() area
integrate(sin_cos_fn, 0, 1)

## 0.9652009 with absolute error < 1.9e-10

# plot function
ggplot(tibble(x = c(0, 0.2, 0.4, 0.6, 0.8, 1.0)), aes(x)) + stat_function(fun = sin_cos_fn)
```



Now with Monte Carlo Integration

```
# run 10^4 uniform samples through our function
unif_rvs <- sin_cos_fn(runif(10^4))

# get estimate of integral area
estimated_integral <- cumsum(unif_rvs) / (1:10^4)

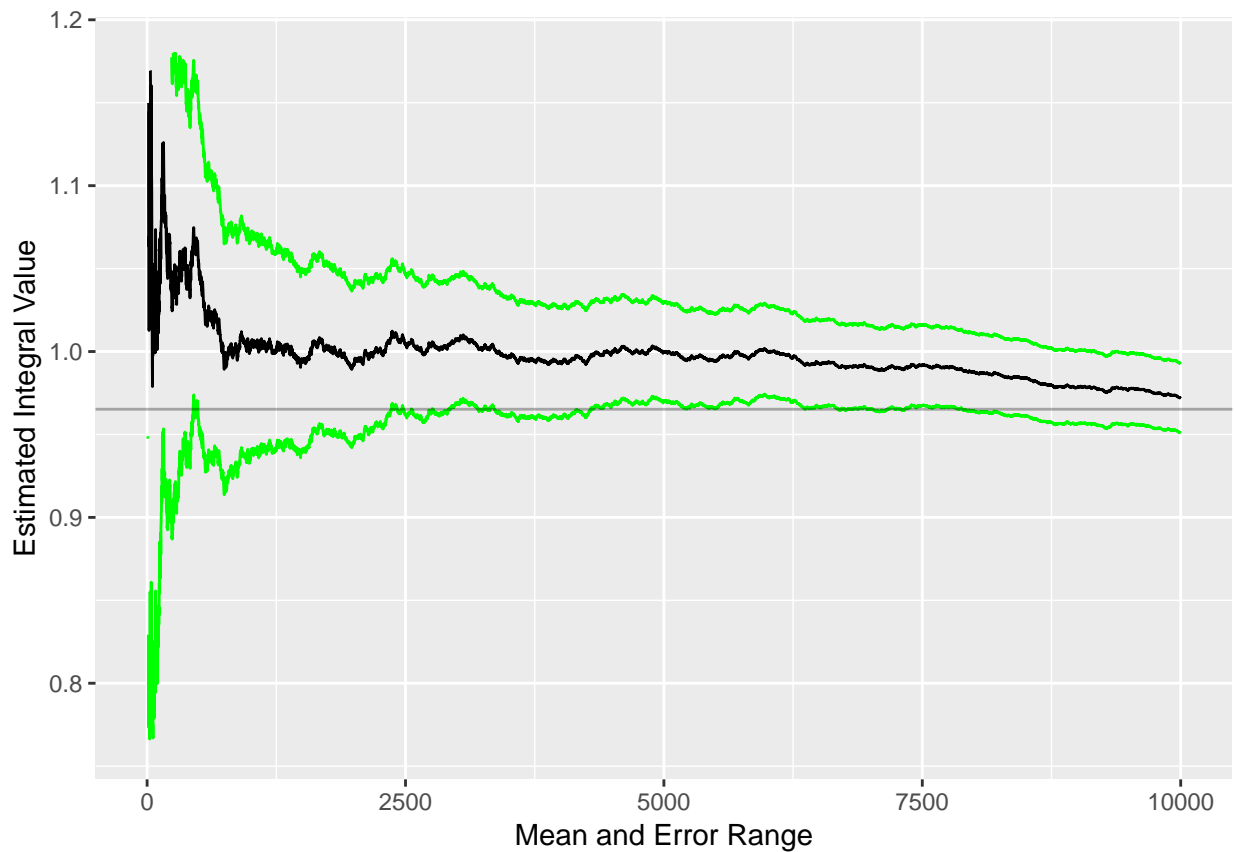
# get estimate of error
estimation_error <- sqrt(cumsum((unif_rvs - estimated_integral)^2)) / (1:10^4)

# add to tibble
estimated <- tibble("Integral" = estimated_integral, "Error" = estimation_error) %>% mutate("Index" = row_number())

# value of integral
max(estimated_integral[9750:10^4])

## [1] 0.9751649

# plot
ggplot(estimated, aes(x = estimated$Index, y = estimated$Integral)) + geom_line() +
  xlab("Mean and Error Range") + ylab("Estimated Integral Value") +
  ylim(mean(unif_rvs) + 20 * c(-estimated$error[10^4], estimated$error[10^4])) +
  geom_line(aes(y = estimated$Integral + 2 * estimated$error), color = "green") +
  geom_line(aes(y = estimated$Integral - 2 * estimated$error), color = "green") +
  geom_hline(yintercept = 0.9652009, alpha = 0.3)
```



While the bonus brought by the simultaneous evaluation of the error of the monte carlo estimate can not be

disputed, we must be aware that it is only trustworthy as far as v_n is a proper estimate of the variance of \bar{h}_n . In critical situations where v_n does not converge at all or does not even converge fast enough for a CLT to apply, this estimate and the confidence region associated with it can not be trusted.

Example 3.4

Given a normal $\mathcal{N}(0, 1)$ sample of size n , (x_1, \dots, x_n) , the approximation of

$$\Phi(t) = \int_{-\infty}^t \frac{1}{\sqrt{2\pi}} e^{-\frac{y^2}{2}} dy$$

by the Monte Carlo method is

$$\hat{\Phi}(t) = \frac{1}{n} \sum_{i=1}^n \mathbb{I}_{x_i \leq t}$$

with exact variance $\frac{\Phi(t)[1-\Phi(t)]}{n}$, since the variables $\mathbb{I}_{x_i \leq t}$ are independent Bernoulli random variables with success probability $\Phi(t)$.

3.3 | Importance Sampling